H2O Driverless AI

Webinar

H2O.ai, Inc, Jul 27 2017

Why Driverless?

- Recipes for Problem Solving
- Automatic Feature Engineering
- Fast GPU Editions of ML and DL -- Allows for automation.
- Model Interpretation
- Automatic Visualization
- Who needs Driverless AI?

H2O.ai

Shortage of Data Scientists

"The United States alone faces a shortage of 140,000 to 190,000 people with analytical expertise and 1.5 million managers and analysts"

–McKinsey Prediction for 2018

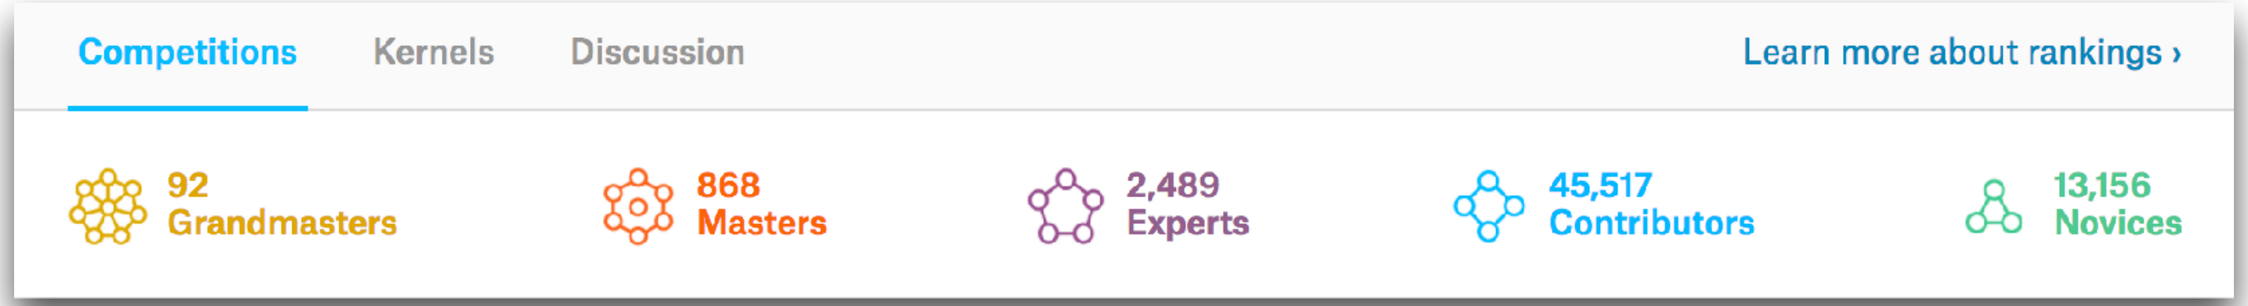

H2O.ai

Typical Enterprise Machine Learning Workflow

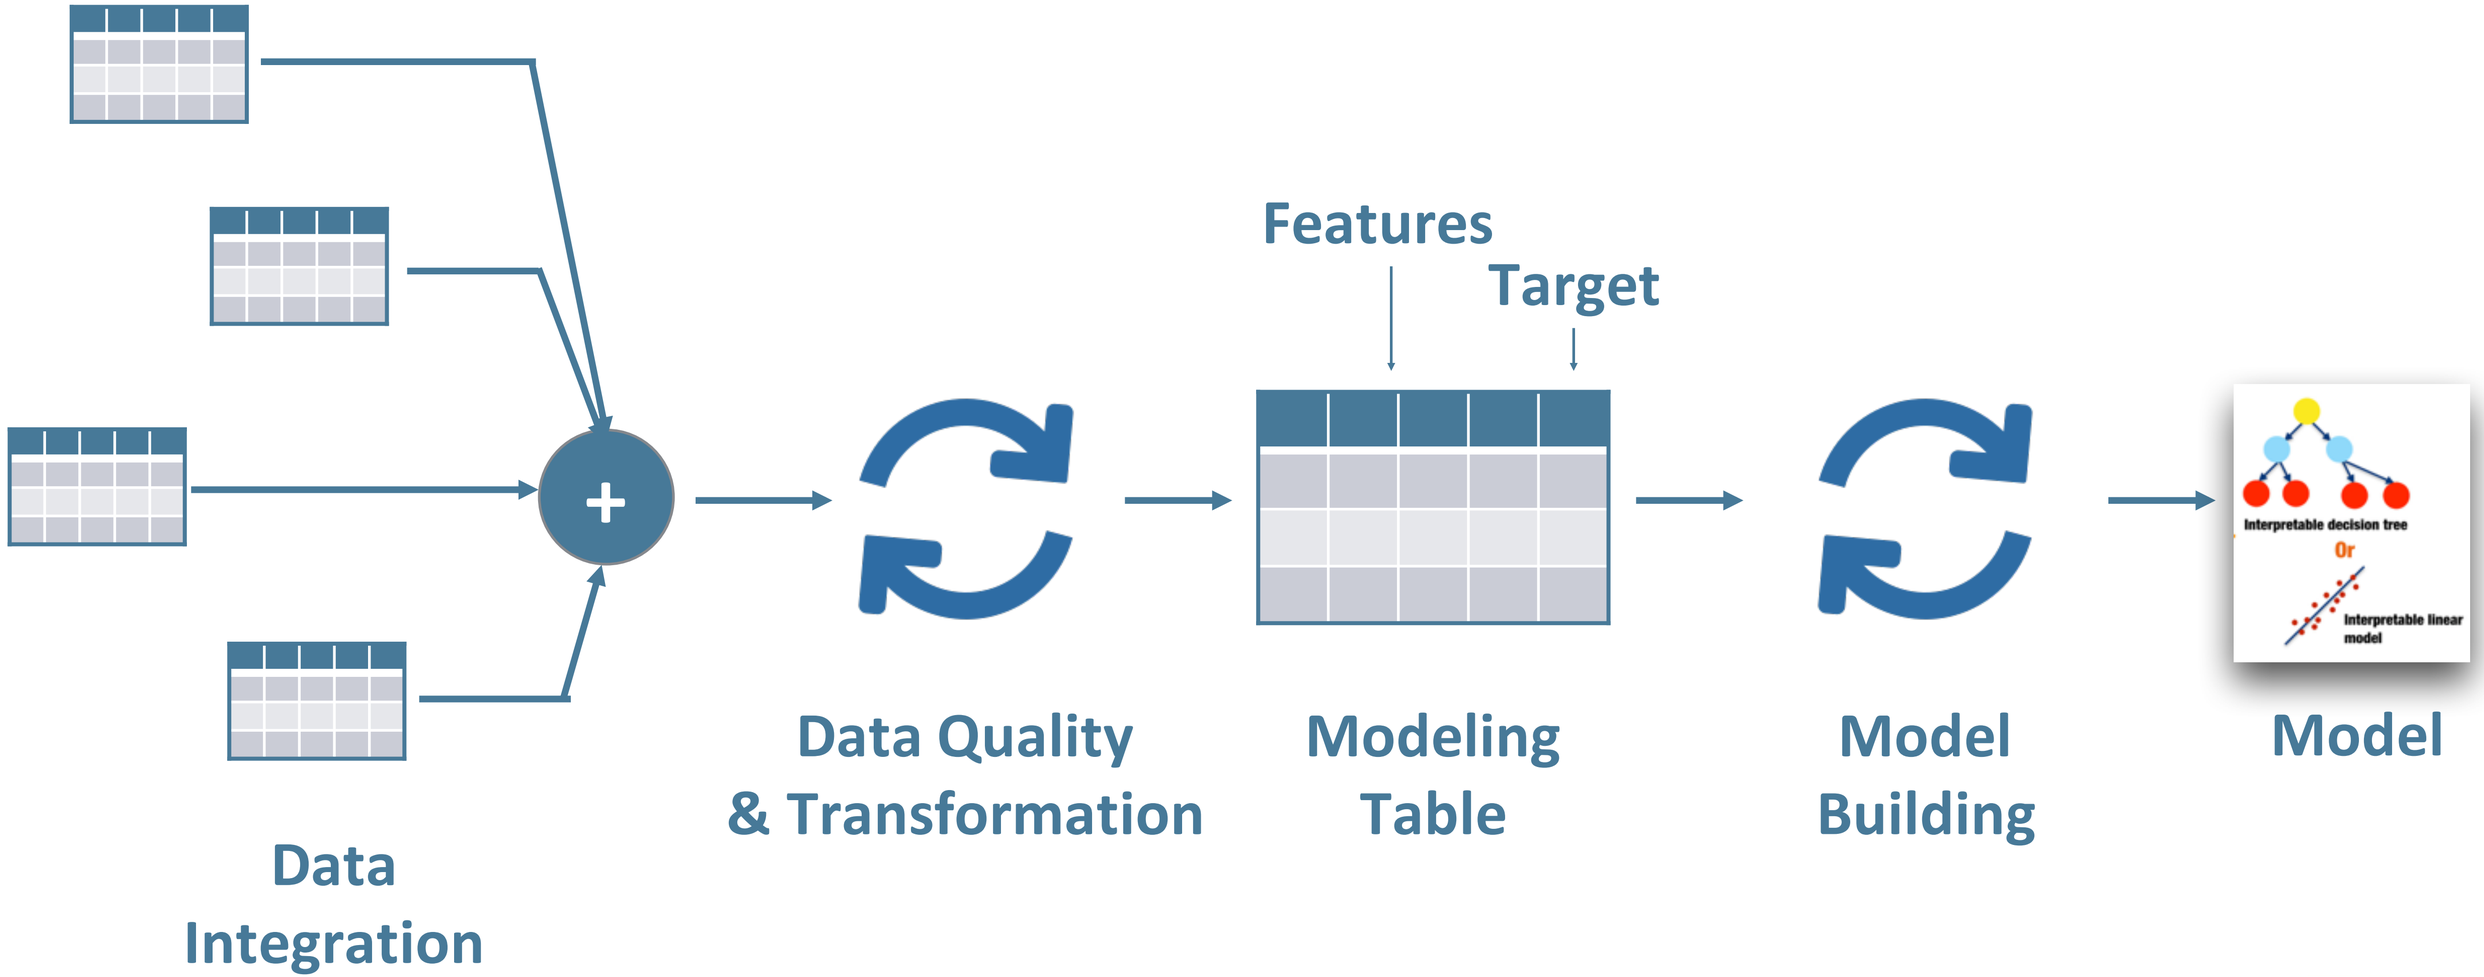

Typical Enterprise Machine Learning Workflow

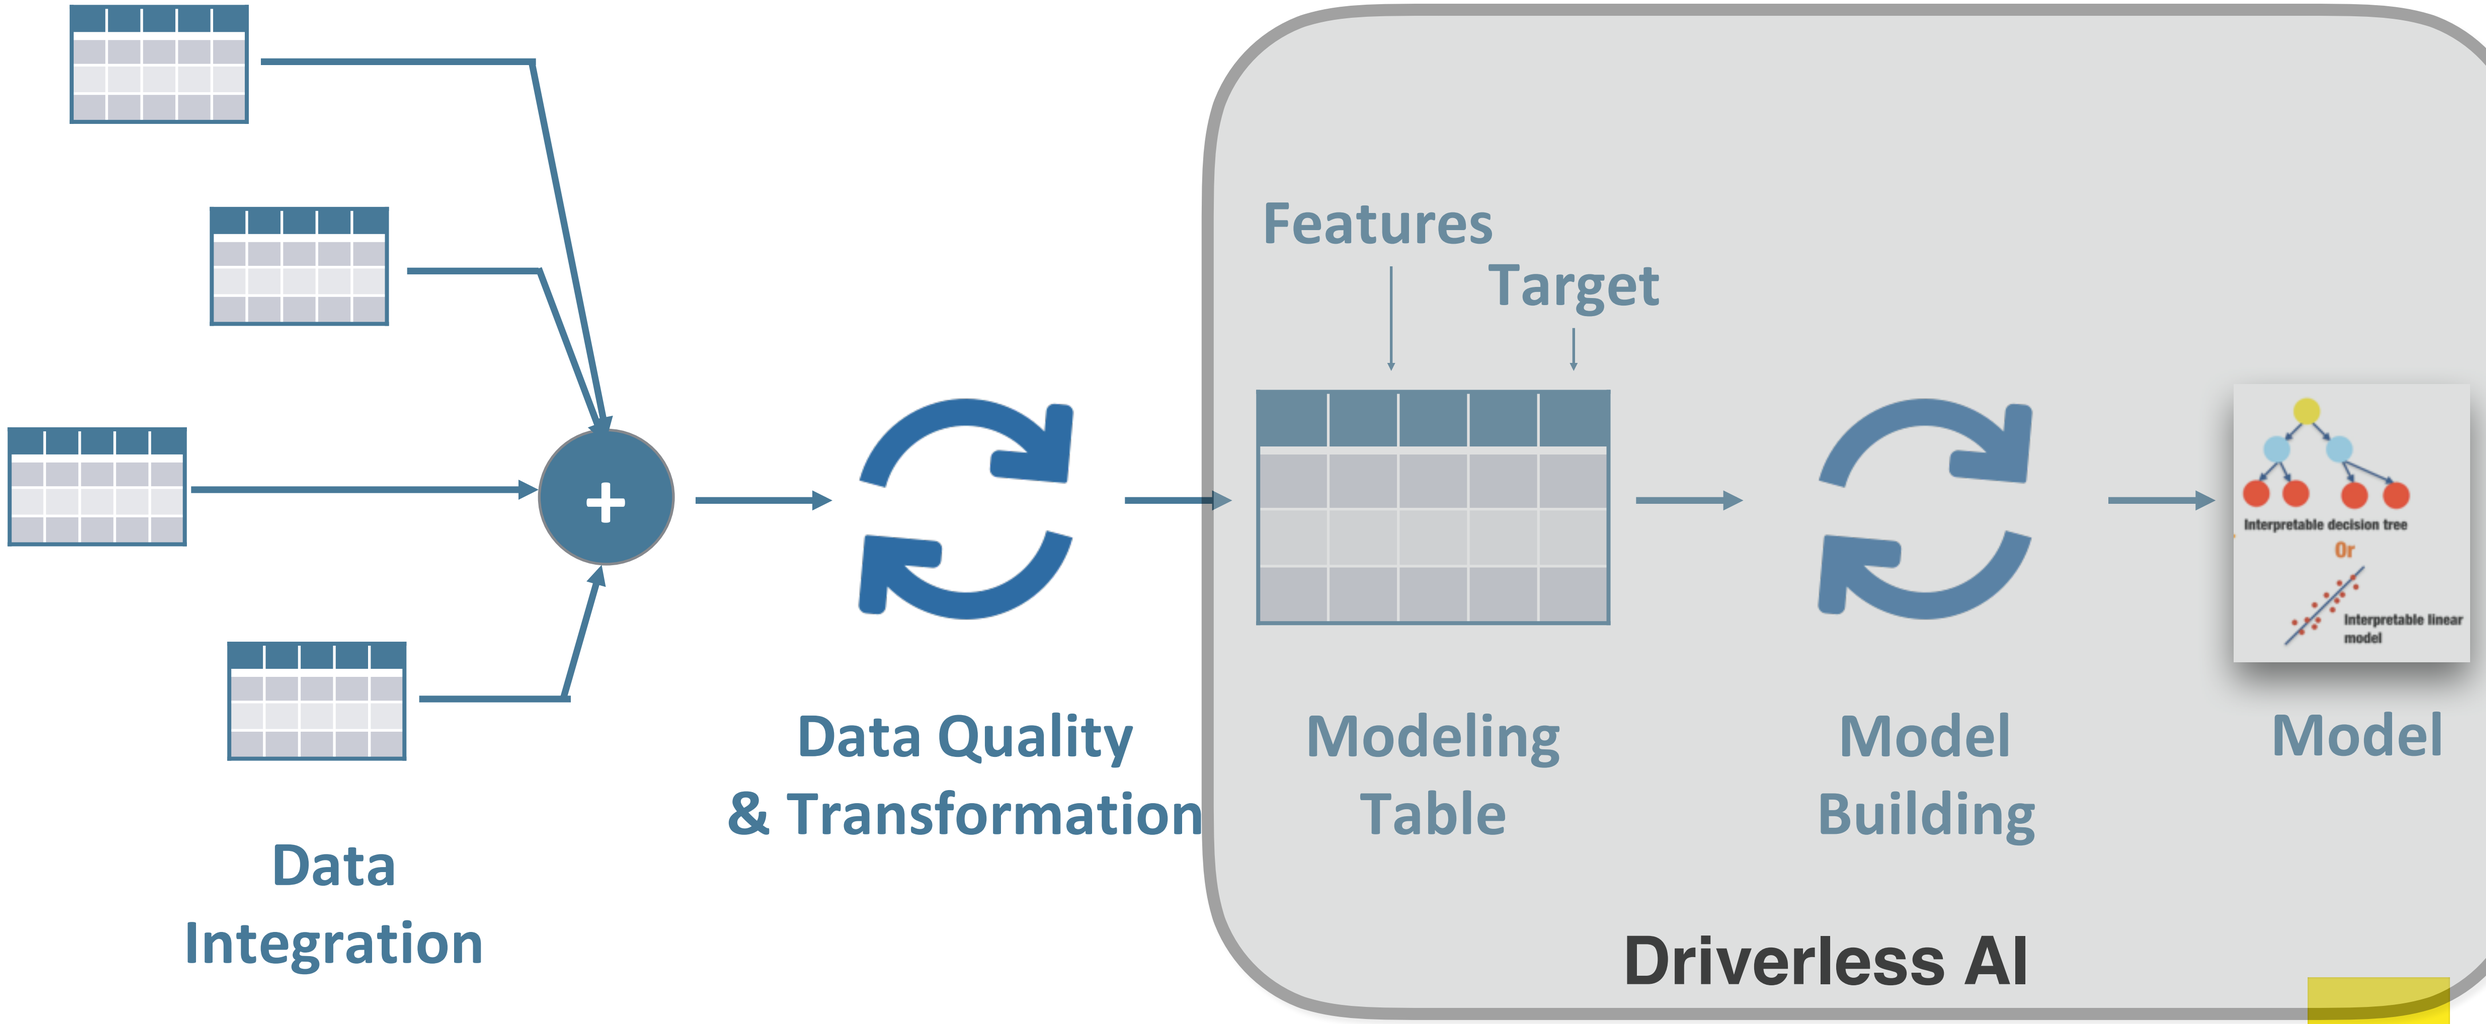

Driverless AI

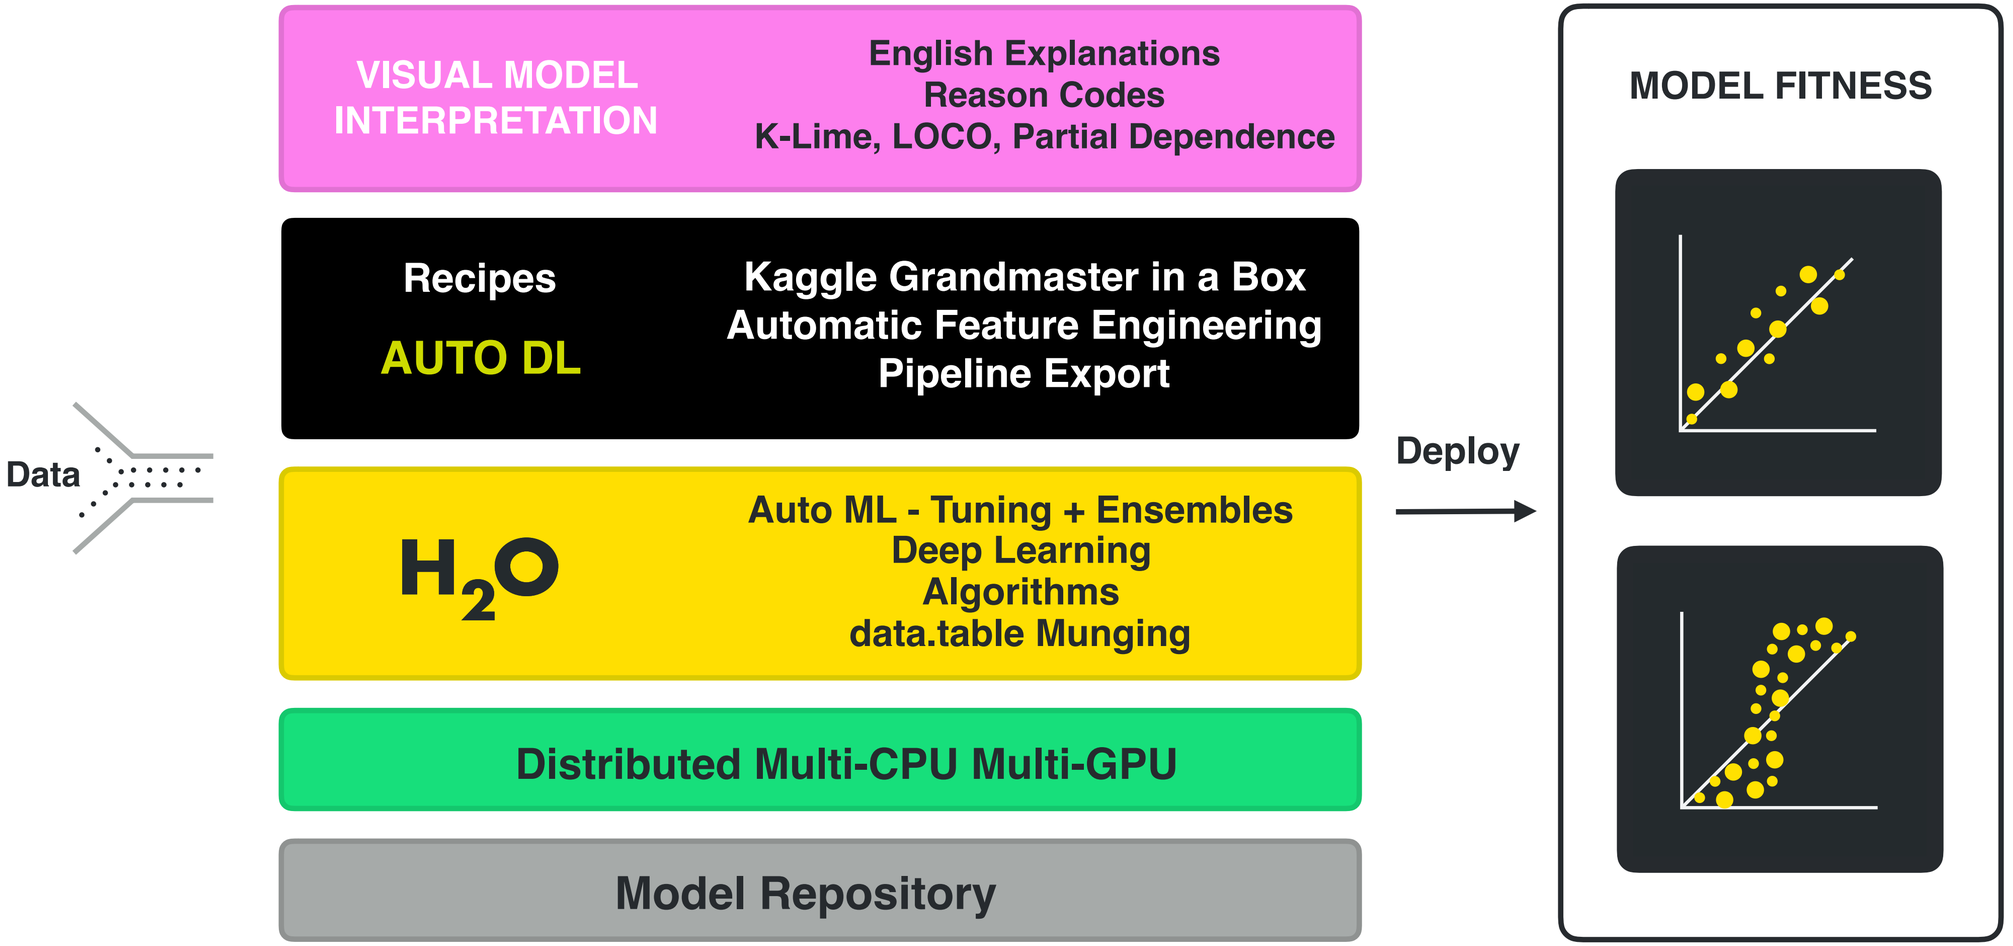

3 Pillars

Speed

Accuracy

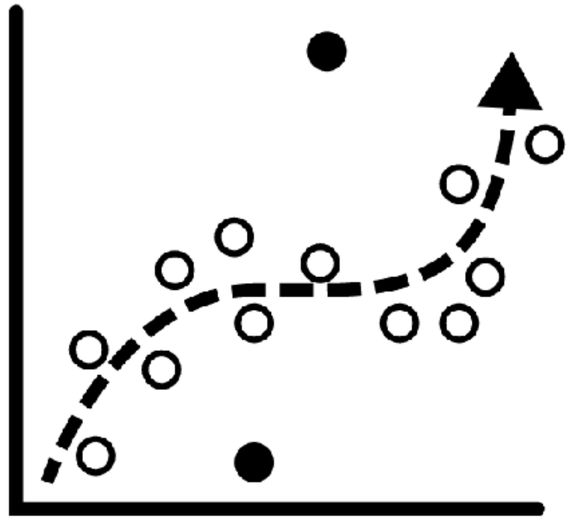

Interpretability

Speed

- GPU acceleration to achieve up to 40x speedups vs CPU
- Multi GPU - XGBoost, GLM, K-Means and more
- Achieve best performance in shortest time

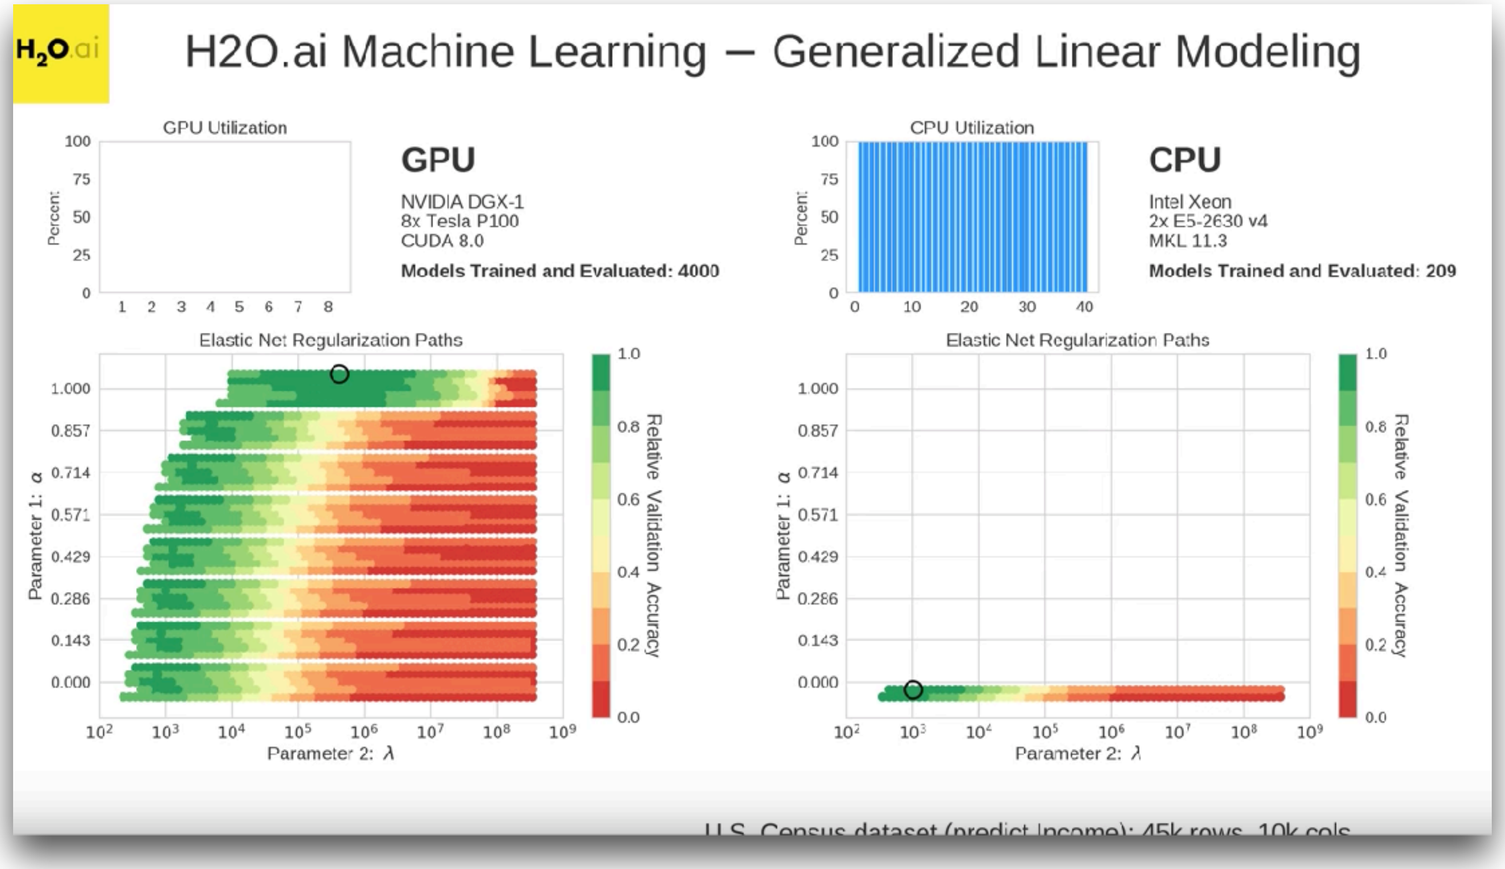

H2O.ai Machine Learning – Generalized Linear Modeling

Time to Train and Evaluate 4000 Models

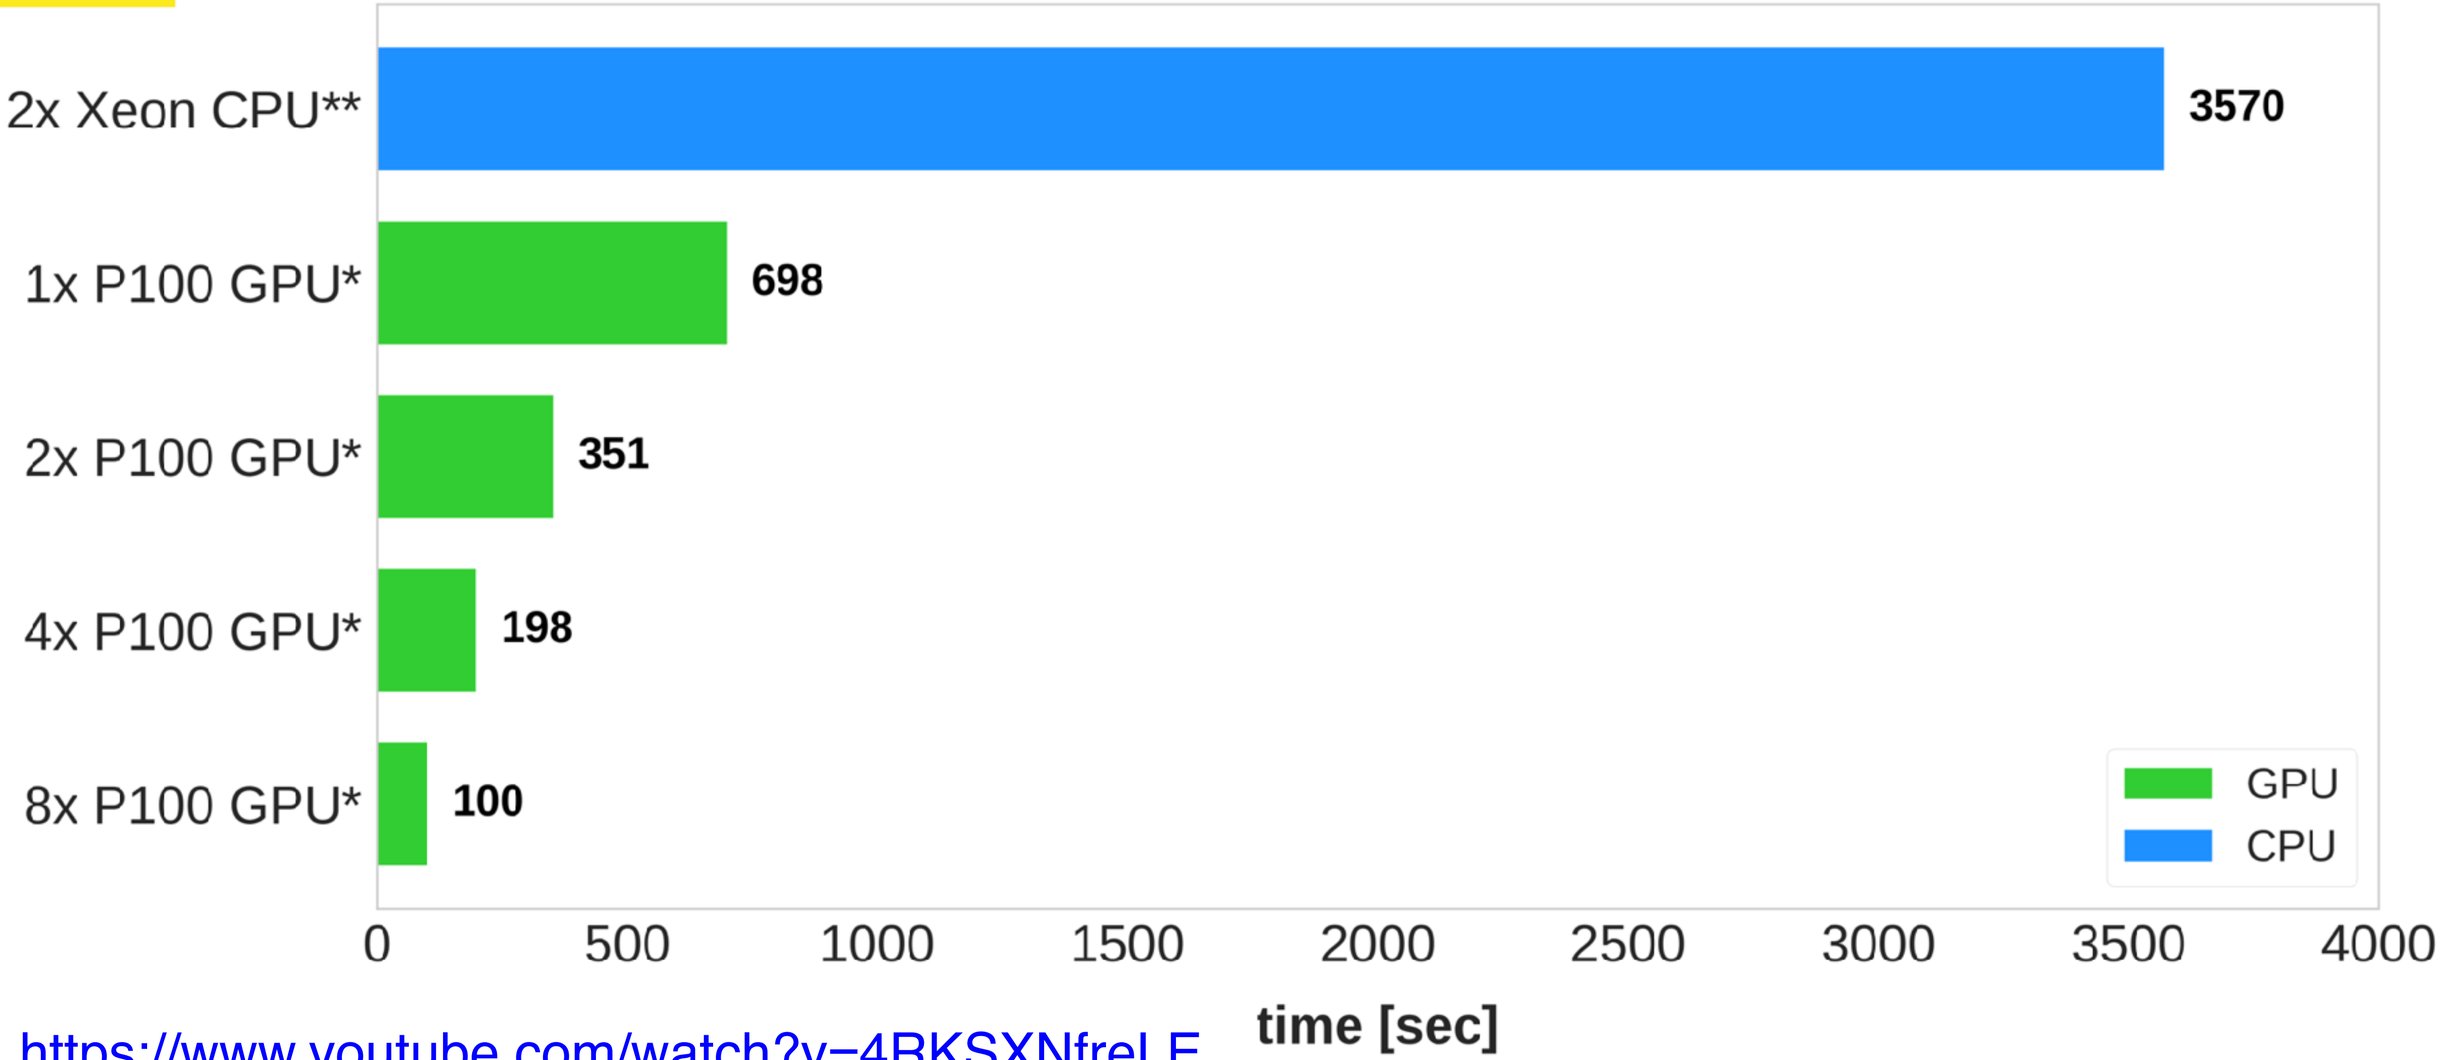

https://www.youtube.com/watch?v=4RKSXNfreLE

http://github.com/h2oai/perf/

*NVIDIA DGX-1, **Dual Intel Xeon E5-2630 v4
U.S. Census dataset (predict Income): 45k rows, 10k cols
Elastic Net Model Parameters: 5-fold cross-validation, $\alpha = \{\frac{i}{7}, i = 0 \ldots 7\}$, full λ-search

H2O.ai Machine Learning – Gradient Boosting Machine

Time to Train 16 H2O XGBoost Models

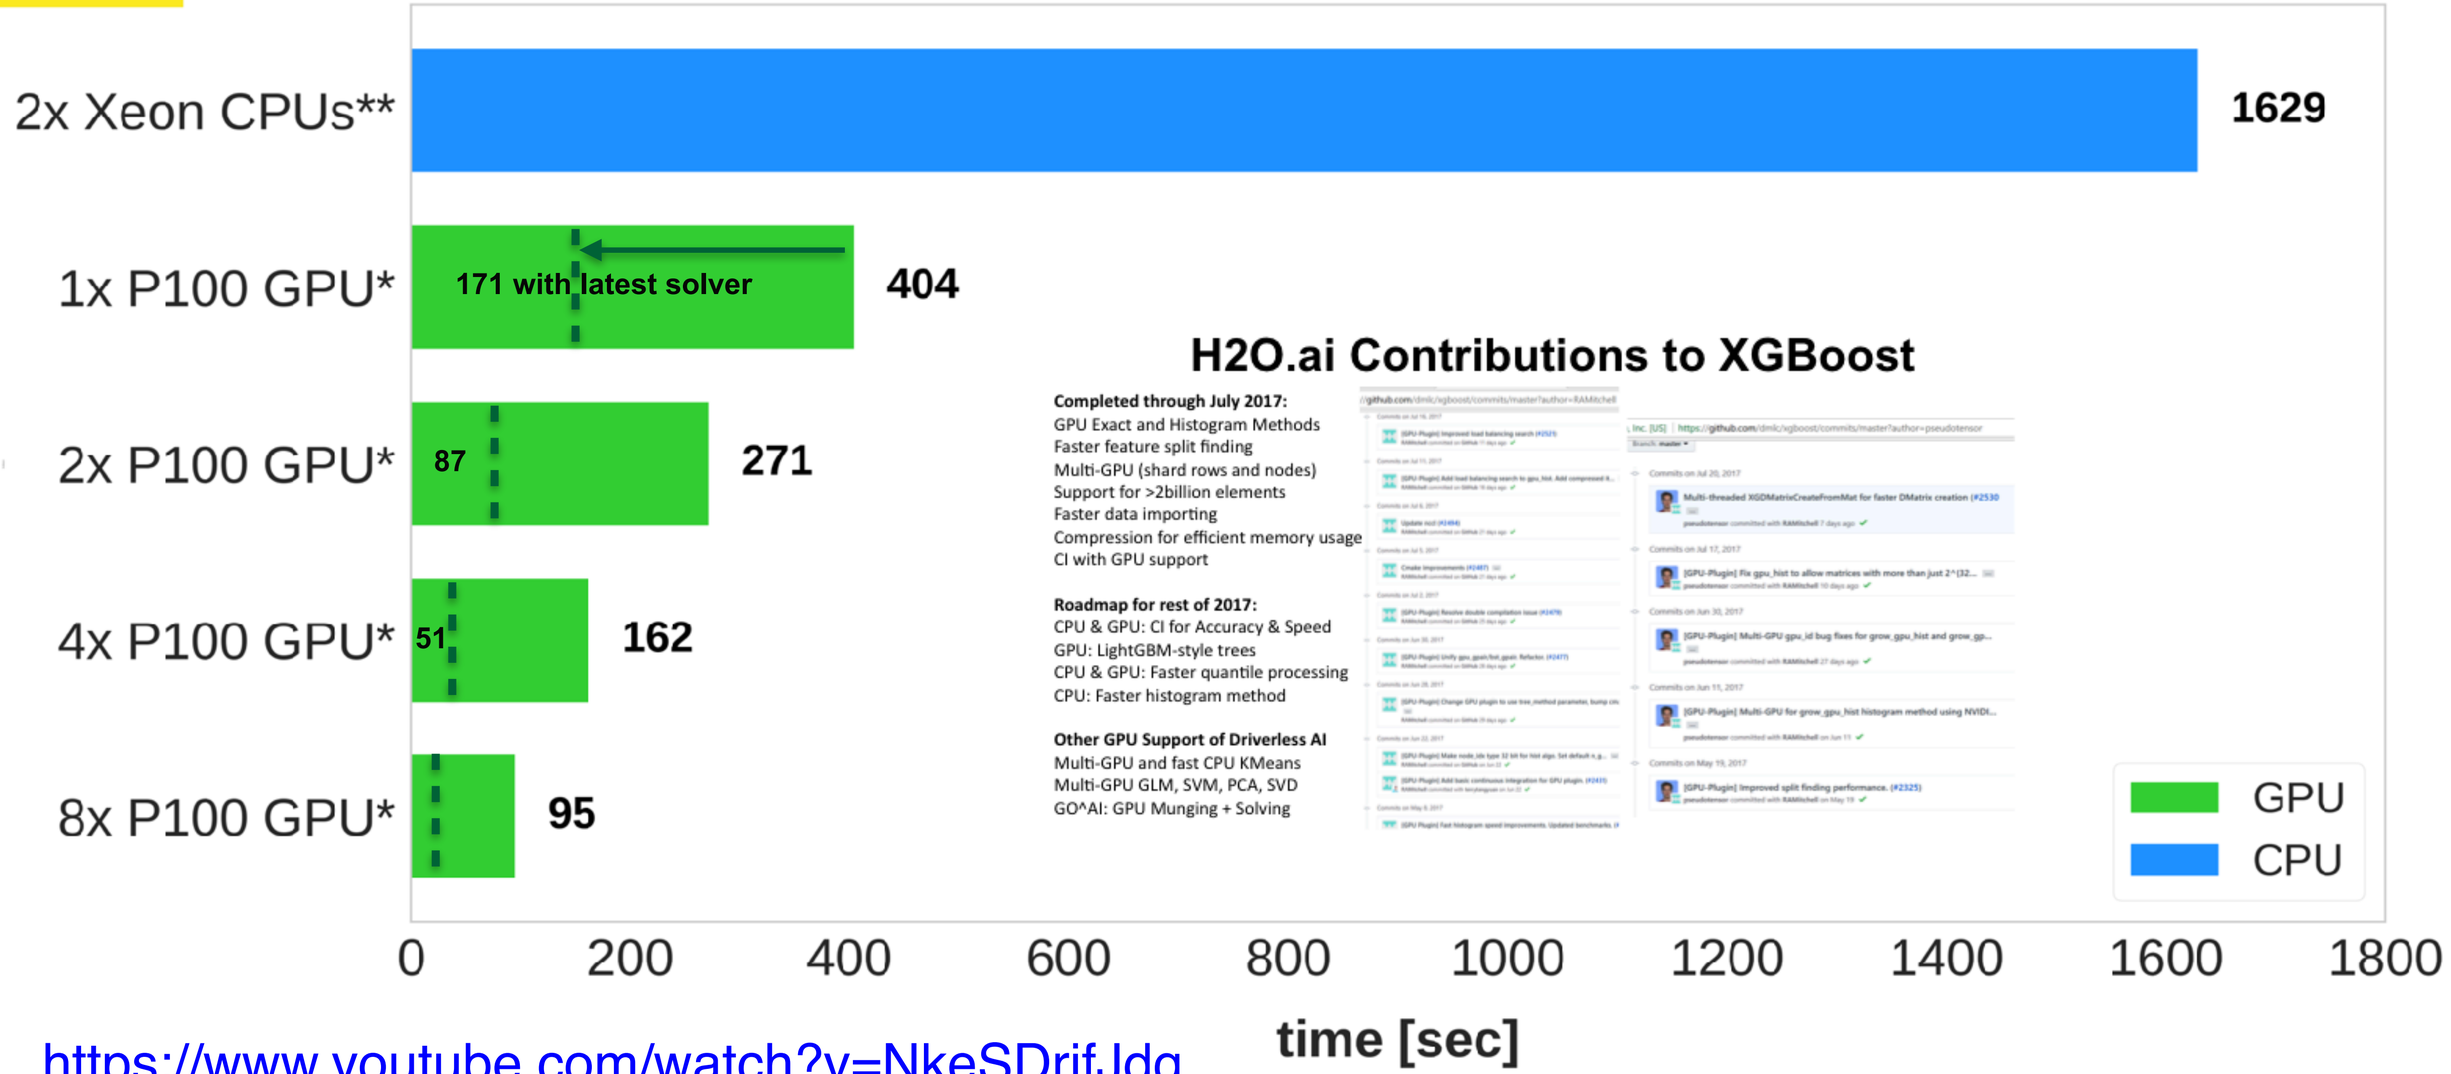

H2O.ai Contributions to XGBoost

Completed through July 2017:
GPU Exact and Histogram Methods
Faster feature split finding
Multi-GPU (shard rows and nodes)
Support for >2billion elements
Faster data importing
Compression for efficient memory usage
CI with GPU support

Roadmap for rest of 2017:
CPU & GPU: CI for Accuracy & Speed
GPU: LightGBM-style trees
CPU & GPU: Faster quantile processing
CPU: Faster histogram method

Other GPU Support of Driverless AI
Multi-GPU and fast CPU KMeans
Multi-GPU GLM, SVM, PCA, SVD
GO^AI: GPU Munging + Solving

https://www.youtube.com/watch?v=NkeSDrifJdg

*NVIDIA DGX-1, **Dual Intel Xeon E5-2630 v4

http://github.com/h2oai/perf/

Higgs dataset (binary classification): 1M rows, 29 cols; max_depth: {6,8,10,12}, sample_rate: {0.7,0.8,0.9,1.0}

H2O.ai Machine Learning – k-Means Clustering

Time to run 1000 Lloyds iterations for k=1000 clusters

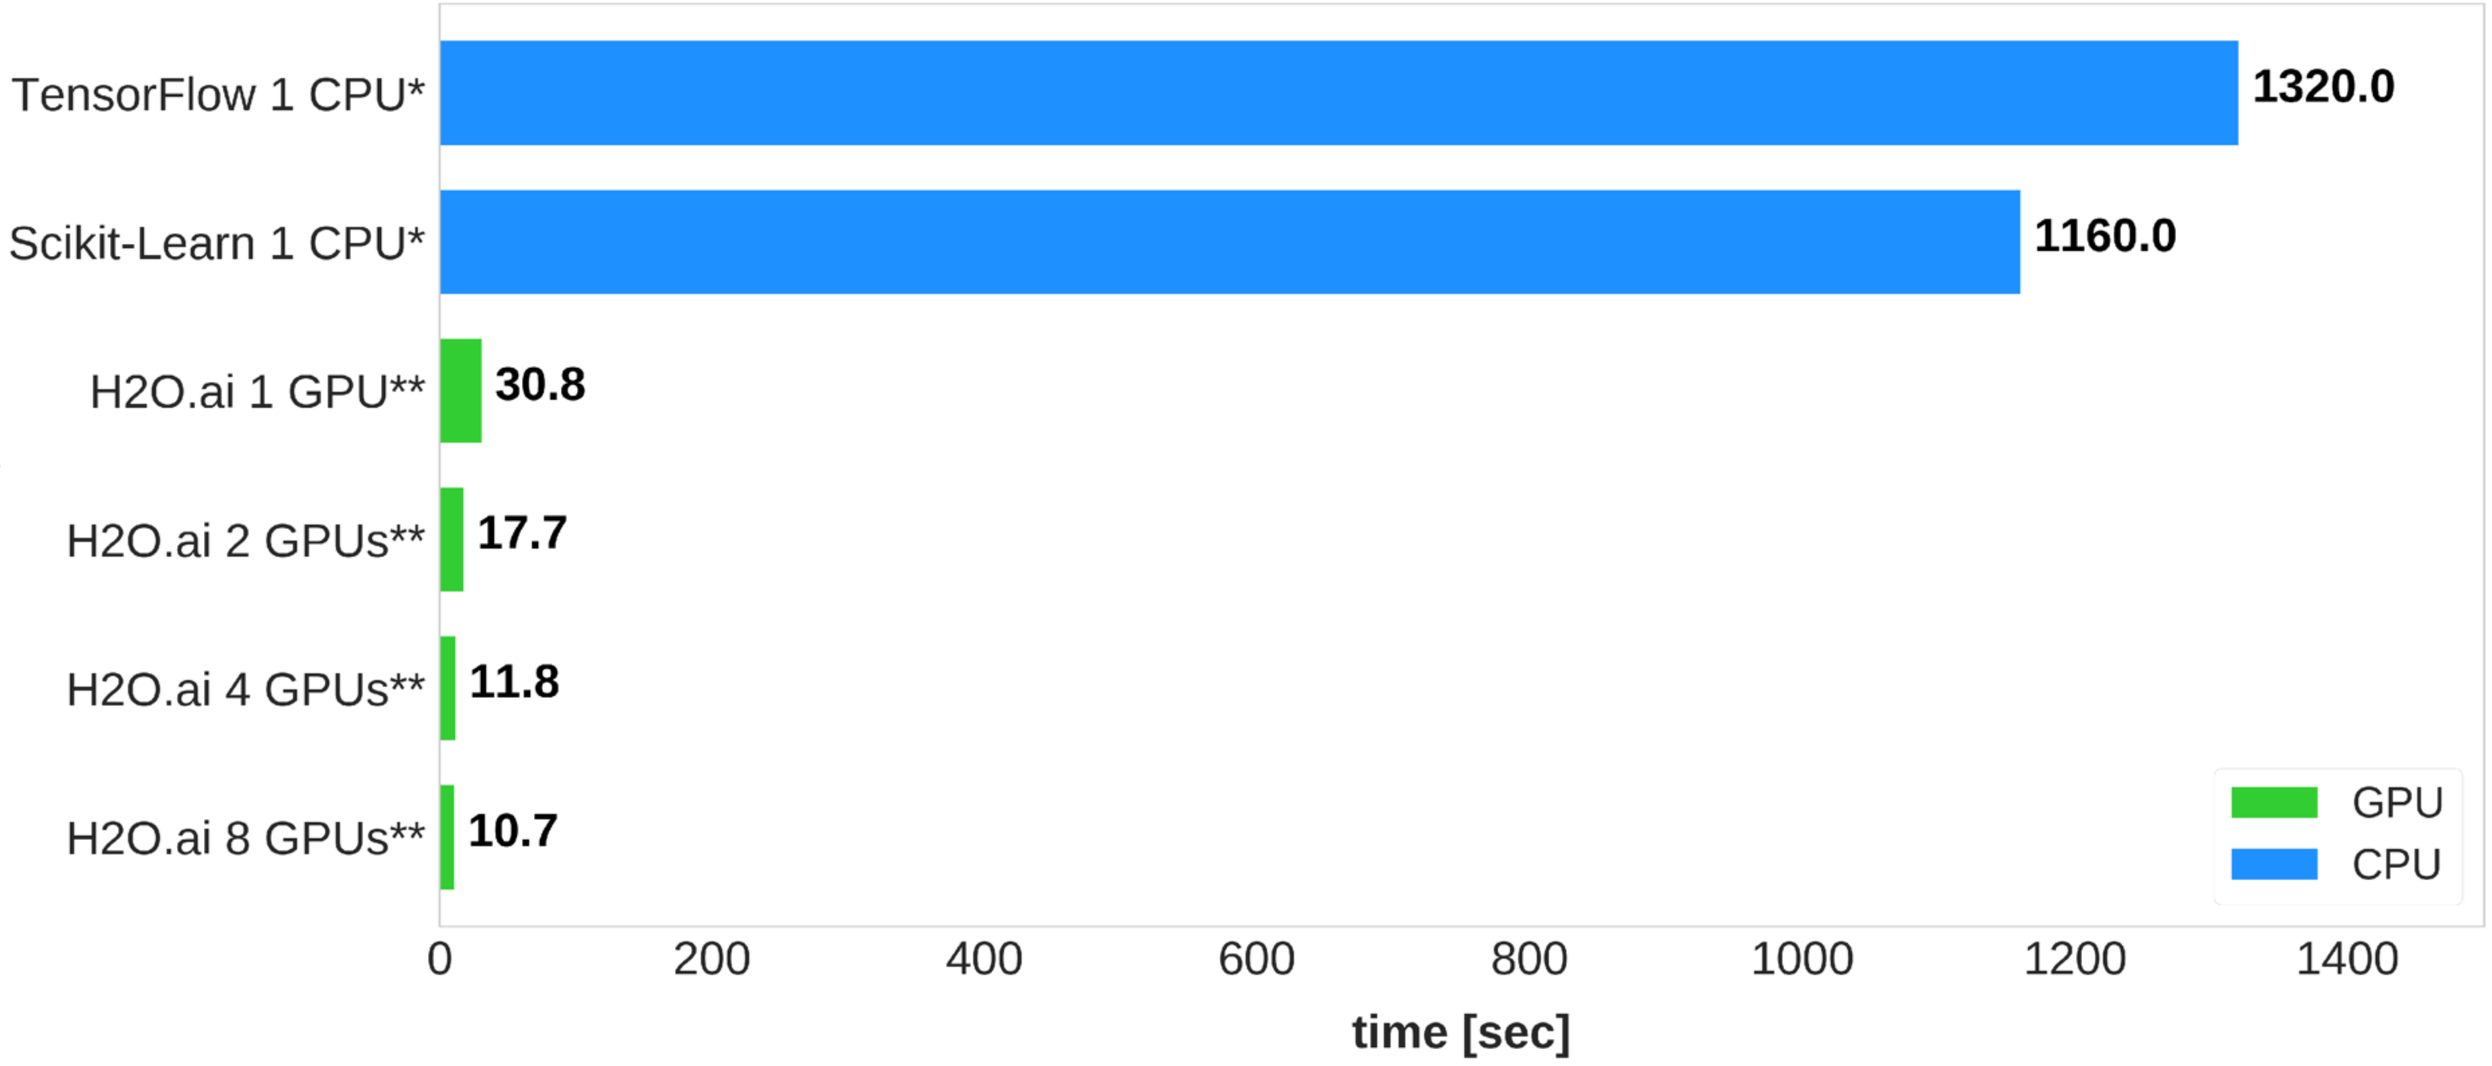

Kaggle Homesite Home Insurance Claims Predictions Dataset (261k rows, 298 cols)
k-Means Clustering (Lloyds), random initialization, 1000 centroids, 1000 iterations
Hardware: *Intel i7 5820k (6-core), **NVIDIA Tesla P100 (DGX-1)

http://github.com/h2oai/perf/

Accuracy

- Automatic feature engineering to increase accuracy - AlphaGo for AI
- Automatic Kaggle Grandmaster recipes in a box for solving wide variety of use-cases
- Automatic machine learning to find and tune the right ensemble of models

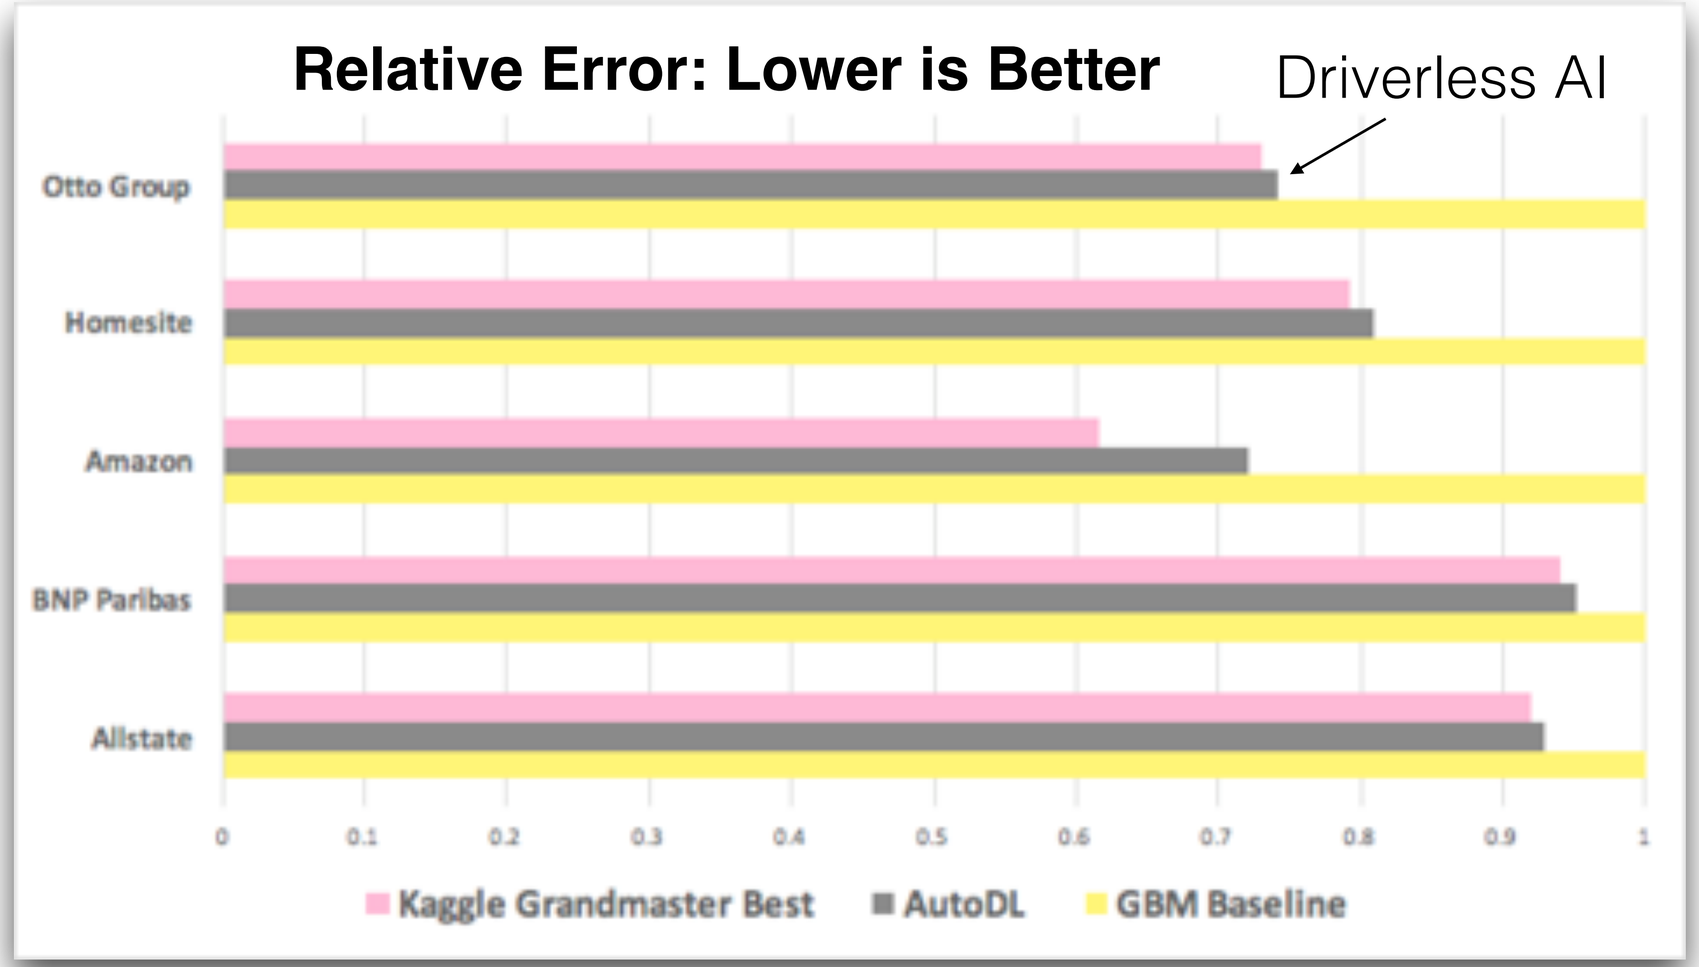

Preliminary results - untuned, single model

What's the "Secret" to Data Science?

"Coming up with features is difficult, time-consuming, requires expert knowledge. "Applied machine learning" is basically feature engineering."

–Andrew Ng

"… some machine learning projects succeed and some fail. What makes the difference? Easily the most important factor is the features used."

–Pedro Domingos

"Good data preparation and feature engineering is integral to better predictions."

–Kazanova, H2O.ai Kaggle Grandmaster #2 (former #1)

H2O.ai

Why is it so difficult to be a good Data Scientist?

Data matters: need access to business-relevant data and need domain knowledge about how features interact with each other

Data Science Recipes

Powerful feature transformations (like target encoding) can lead to overfit models if done wrong, need strong math & statistics skills

Need to run thousands of experiments to reach robust conclusions, need computer science skills and access to compute hardware

GPUs, Automation

H2O.ai

Interpretability

- Interpretability for debugging, not just for regulators
- Get reason codes and model interpretability in plain english
- K-Lime, LOCO, partial dependence and more

Automatic Visualization

H2O.ai

CLUMPY SCATTERPLOTS

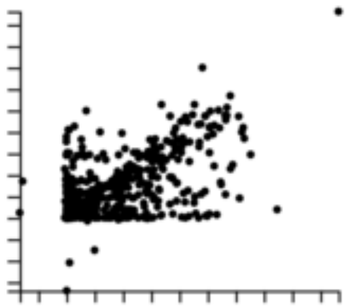

MONOTONIC SCATTERPLOTS

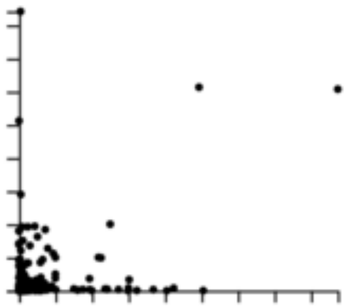

SPIKEY HISTOGRAMS

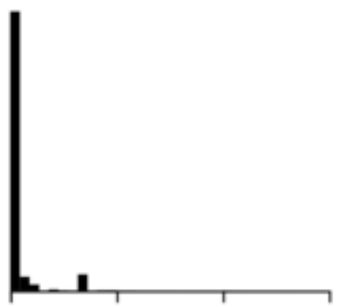

OUTLIERS

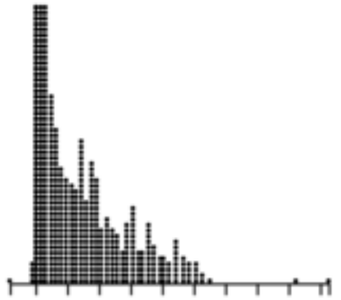

HEATMAPS

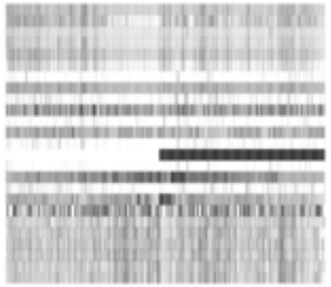

H2O.ai

Next: Live Demo!

Installation and Start

```
# Load the Driverless AI docker image
docker load < downloaded-h2oai-driverless-ai-image.tar
```

```
# Start the Driverless AI docker image
nvidia-docker run --rm -u `id -u`:`id -g` -p 12345:12345 -v `pwd`/data:/data -v `pwd`/
 ↪log:/log opsh2oai/h2oai-runtime
```

The Team Behind Driverless AI


Marios Michailidis


Dmitry Larko


Branden Murray


Mark Landry

AutoDL


Matt Dowle


Pasha Stetsenko

Datatable


Patrick Hall


Mark Chan


Navdeep Gill

MLI


Arno Candel


Jon Mckinney

Rory Mitchell

GPU Acceleration


Leland Wilkinson


Prithvi Prabhu


Justin Loyola

AutoViz

H2O.ai

Upcoming Webinars

- Machine Learning Interpretability - Patrick Hall - **Aug 17**
- AutoDL - Dmitry Larko - **Aug 24**
- Automatic Visualization - **Sep (TBD)**

H2O.ai