



An Introduction to CUDA Programming

S7699 – Session 1 of 4



Chris Mason

Product Manager, Acceleware

GPU Technology Conference

Date: May 8, 2017

About Acceleware

Programmer Training

- CUDA and other HPC training classes
- Over 100 courses taught
- <http://acceleware.com/training>

Consulting Services

- Projects for Oil & Gas, Medical, Finance, Security and Defence, CAD, Media & Entertainment
- Mentoring, code review and complete project implementation
- <http://acceleware.com/services>

GPU Accelerated Software

- Seismic imaging & modeling
- Electromagnetics



Seismic Imaging & Modeling

AxWAVE™

- Seismic forward modeling
- 2D, 3D, constant and variable density models
- High fidelity finite-difference modeling

AxRTM™

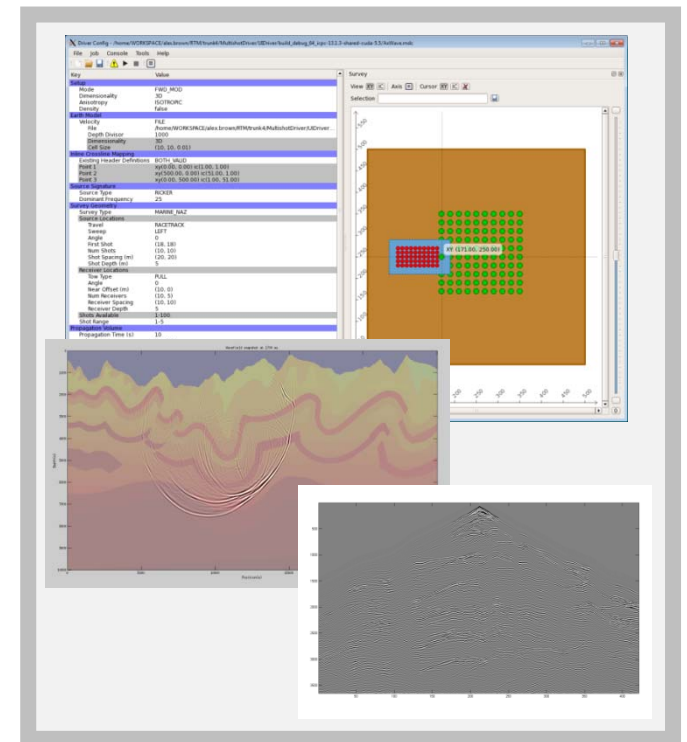
- High performance Reverse Time Migration application
- Isotropic, VTI and TTI media

AxFWI™

- Inversion of the full seismic data to provide an accurate subsurface velocity model
- Customizable for specific workflows

HPC Implementation

- Optimized for NVIDIA Tesla GPUs
- Efficient multi-GPU scaling




Electromagnetics

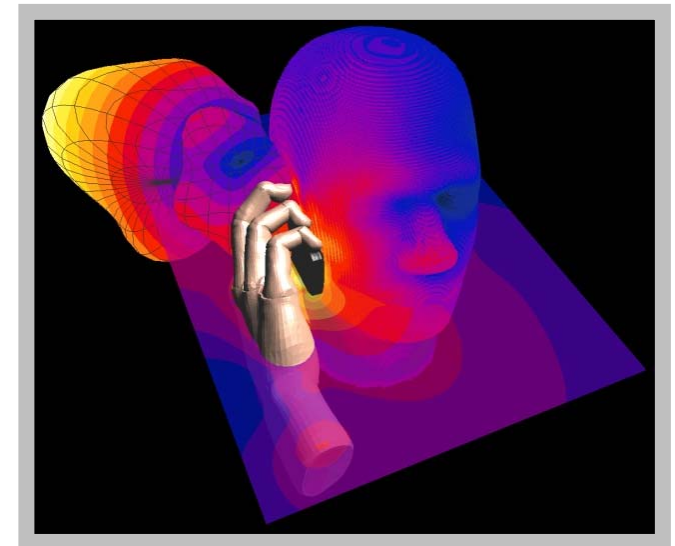
AxFDTD™

- Finite-Difference Time-Domain Electromagnetic Solver
- Optimized for NVIDIA GPUs
- Sub-gridding and large feature coverage
- Multi-GPU, GPU clusters, GPU targeting

Available from:



Agilent Technologies



Consulting Services

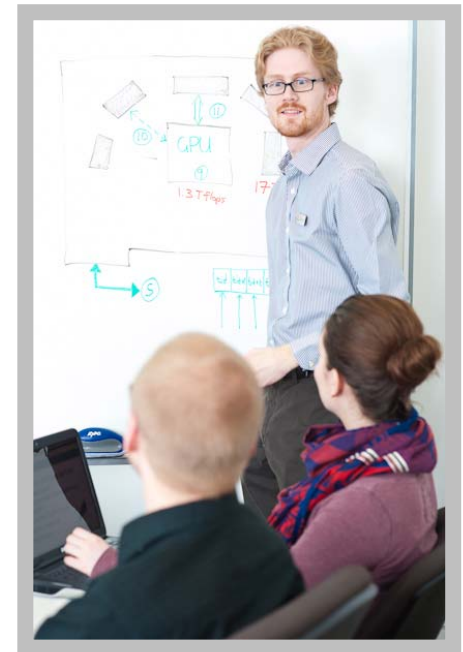
Industry	Application	Work Completed	Results
Finance	Option Pricing	Debugged & optimized existing CUDA code Implemented the Leisen-Reimer version of the binomial model for stock option pricing	30-50x performance improvement compared to single-threaded CPU code
Security & Defense	Detection System	Replaced legacy Cell-based infrastructure with GPUs Implemented a GPU accelerated X-ray iterative image reconstruction and explosive detection algorithms	Surpassed the performance targets Reduced hardware cost by a factor of 10
CAE	SIMULIA Abaqus	Developed a GPU accelerated version Conducted a finite-element analysis and developed a library to offload LDLT factorization portion of the multi-frontal solver to GPUs	Delivered an accelerated (2-3x) solution that supports NVIDIA and AMD GPUs
Medical	CT Reconstruction Software	Developed a GPU accelerated application for image reconstruction on CT scanners and implemented advanced features including job batch manager, filtering and bad pixel corrections	Accelerated back projection by 31x
Oil & Gas	Seismic Application	Converted MATLAB research code into a standalone application & improved performance via algorithmic optimizations	20-30x speedup

Programmer Training

- CUDA and other HPC training classes
- Public, private onsite, and online courses
- Teachers with real world experience
- Hands-on lab exercises
- Progressive lectures
- Small class sizes to maximize learning
- 90 days post training support

“The level of detail is fantastic. The course did not focus on syntax but rather on how to expertly program for the GPU. I loved the course and I hope that we can get more of our team to take it.”

Jason Gauci, Software Engineer
Lockheed Martin




Outline

CUDA overview

Data-parallelism

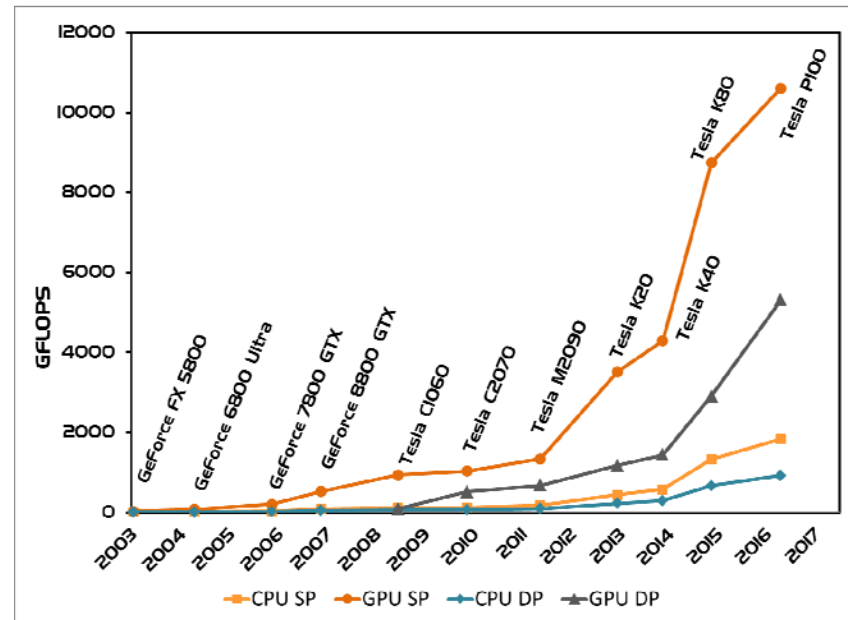
GPU programming model

- GPU kernels
- Host vs. device responsibilities
- CUDA syntax
- Thread hierarchy








Why use GPUs? Performance!

GPU advances are outpacing CPU advances
Continuing Moore's Law?



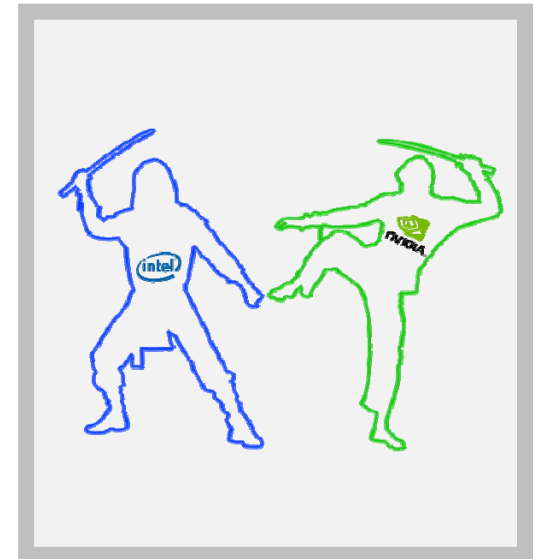
Why use GPUs? Performance!

	Intel Xeon E5-2699v4 (Broadwell-EP) 	NVIDIA Tesla K80 (Kepler) 	NVIDIA Tesla M60 (Kepler) 	NVIDIA Tesla P100 (Pascal) 	NVIDIA Jetson TX2 (Pascal) 
Processing Cores	22	4992	4096	3584	8 ARM + 256 Pascal
Clock Frequency	2.2-3.6GHz	0.562-0.875GHz	0.900-1.180GHz	1.328-1.48GHz	0.854 – 1.465GHz
Memory Bandwidth	76.8 GB/s / socket	480GB/s	320GB/s	720GB/s	58.4GB/s
Peak Tflops (single)	1.83 @ 2.6GHz	8.74 @ 0.875GHz	9.68 @ 1.180GHz	10.6 @ 1.48GHz	0.75@1.465GHz
Peak Tflops (double)	0.915 @ 2.6GHz	2.91 @ 0.875GHz	0.30 @ 1.180GHz	5.3 @ 1.48GHz	0.023@1.465GHz
Gflops/Watt (single)	12.62	29.1	32.2	35.3	50
Total Memory	>>24GB	24GB	16GB	16GB	8GB

GPU Potential Advantages

Tesla P100 vs. Xeon E5-2699 v4

- 5.8x more single-precision floating-point throughput
- 5.8x more double-precision floating-point throughput
- 9.4x higher memory bandwidth

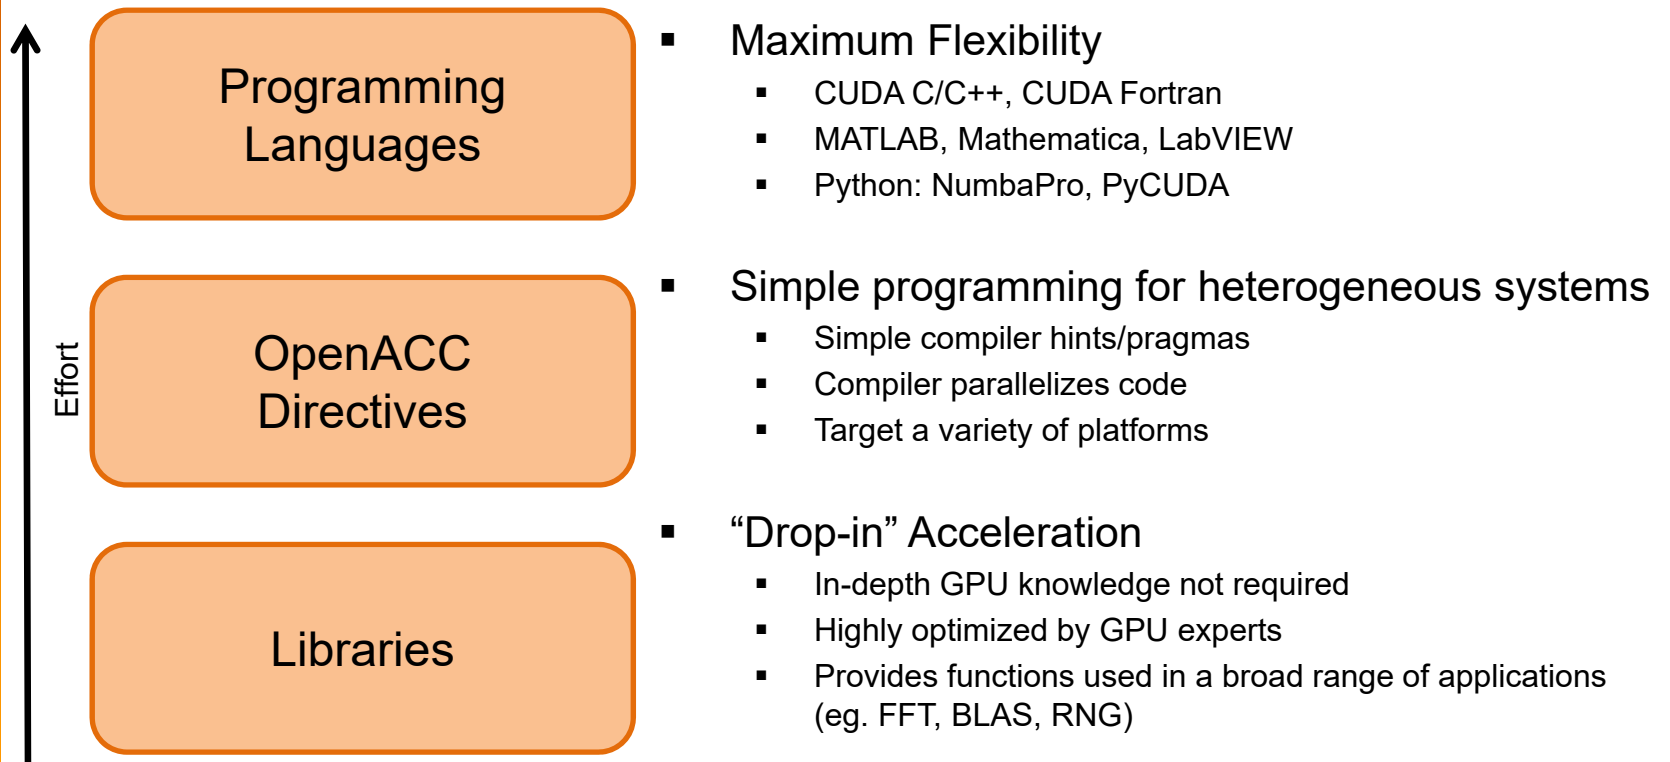


GPU Disadvantages

- Architecture not as flexible as CPU
- Must rewrite algorithms and maintain software in GPU languages
- Discrete GPUs attached to CPU via relatively slow PCIe
 - 32GB/s bi-directional for PCIe 3.0 16x
 - 40GB/s bi-directional for NVLink
- Limited memory (though 8-24GB is reasonable for many applications)



Software Approaches for Acceleration



Compute Capability

Hardware architecture version number

Defined by a major and minor version number

- Major version specifies core architecture type
- Minor version refers to incremental improvements and new features

Architecture	Compute Capability	GPUs	Example Features
Tesla	1.0	GeForce 8800, Tesla C870	Base Functionality
Fermi	2.0	GeForce GTX 480, Tesla C2050	Fast Double Precision, Memory Caches
Kepler	3.0	GeForce GTX 680, Tesla K10	Warp Shuffle Functions
	3.5	GeForce GTX Titan Black, Tesla K40	Dynamic Parallelism
	3.7	Tesla K80	More Registers / Shared Memory
Maxwell	5.0	GeForce GTX 750 Ti, Tegra X1	Power Efficient Architecture
	5.2	Tesla M40, Tesla M60	More Shared Memory
Pascal	6.0	Tesla P100	Half Precision (FP16)
	6.1	Titan Xp	Int8
	6.2	Tegra P1	Int8 + FP16

CUDA Overview

- Compute Unified Device Architecture (CUDA) – Parallel computing platform and programming model architecture developed by NVIDIA
 - Libraries, OpenACC, and programming languages built on top of CUDA
- CUDA C/C++ programming interface consists of:
 - C language extensions to target portions of source code for parallel execution on the device (GPU)
 - A library of C functions that execute on the host (CPU) to interact with the device



Data-Parallel Computing

1. Performs operations on a data set organized into a common structure (eg. an array)
2. A set of **tasks** work collectively and simultaneously on the same structure with each task operating on its own portion of the structure
3. Tasks perform identical operations on their portions of the structure. Operations on each portion must **be data independent!**

Data Dependence

- Data dependence occurs when a program statement refers to the data of a preceding statement

```
a = 2 * x;  
b = 2 * y;  
c = 3 * x;
```

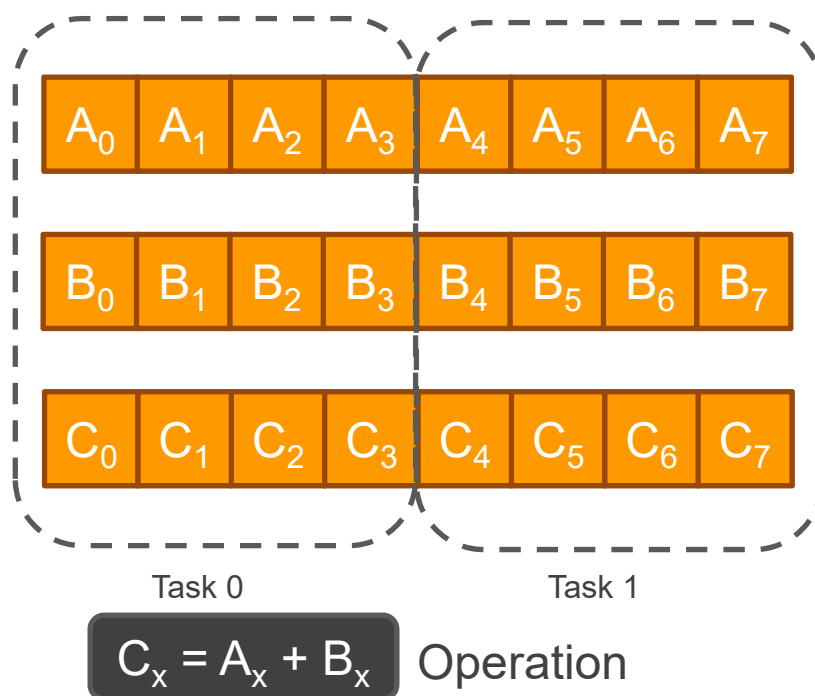
These 3 statements are
independent!

```
a = 2 * x;  
b = 2 * a * a;  
c = b * 9;
```

b depends on *a*, *c* depends
on *b* and *a*!

- Data dependence limits parallelism

Data-Parallel Computing Example

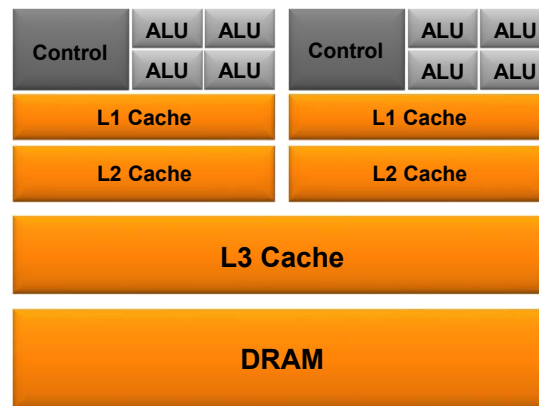


- Data set consisting of arrays A, B, and C
- Same operations performed on each element
 $C_x = A_x + B_x$
- Two tasks operating on a subset of the arrays. Tasks 0 and 1 are independent. Could have more tasks.

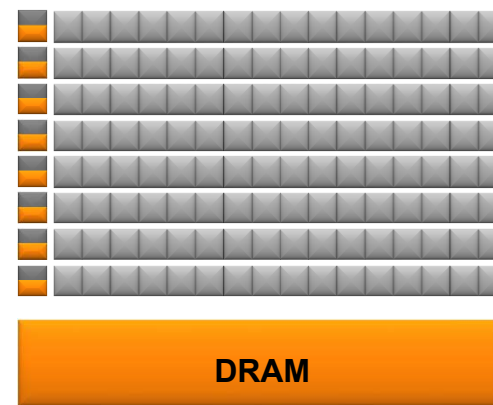
Data-Parallel Computing on GPUs

Data-parallel computing maps well to GPUs:

- Identical operations executed on many data elements in parallel
 - Simplified flow control allows increased ratio of compute logic (ALUs) to control logic




CPU



GPU

The CUDA Programming Model

CUDA is a heterogeneous model, including provisions for both **host** and **device**



The CUDA Programming Model

- Data-parallel portions of an algorithm are executed on the device as **kernels**
 - Kernels are C/C++ functions with some restrictions, and a few language extensions
- Only one kernel is executed at a time
 - Newer GPU architectures relax this restriction
- Each kernel is executed by many **threads**

CUDA Threads

- CUDA threads are conceptually similar to data-parallel tasks
 - Each thread performs the same operations on a subset of a data structure
 - Threads execute independently
- CUDA threads *are not* CPU threads
 - CUDA threads are extremely lightweight
 - Little creation overhead
 - Instant context-switching
- CUDA threads *must* execute the same kernel

CUDA Thread Hierarchy

- CUDA is designed to execute 1000s of threads
- Threads are grouped together into **thread blocks**
- Thread blocks are grouped together into a **grid**


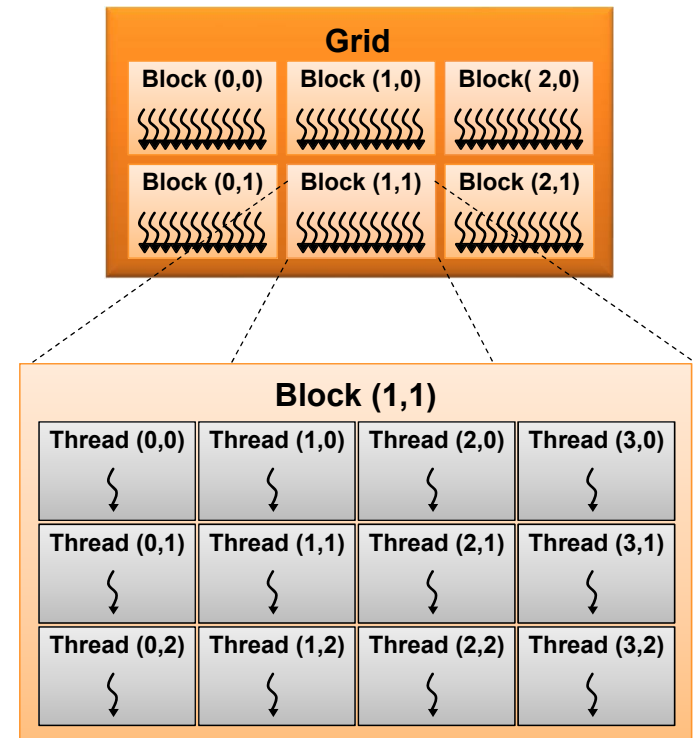


Image courtesy of NVIDIA Corp.

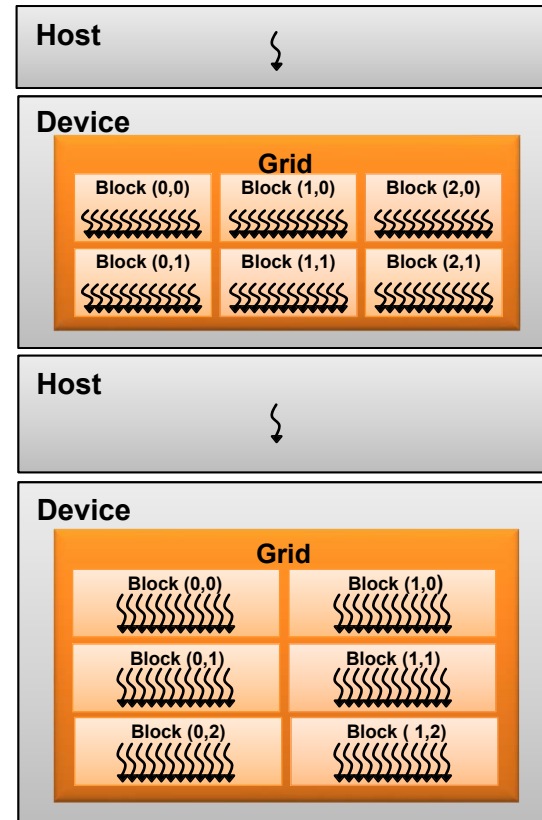
CUDA Thread Hierarchy

- Thread blocks and Grids can be 1D, 2D or 3D
- Dimensions set at launch time
- Thread blocks and grids do not need to have the same dimensionality
 - ie. 1D Grid of 2D Thread Blocks



The CUDA Programming Model

- The host launches kernels
- The host executes serial code between device kernel launches
 - Memory management
 - Data exchange to/from device
 - Error handling



CUDA APIs

Can use CUDA through CUDA C (Runtime API) and low-level Driver API

- **The tutorial presentations use CUDA C**
 - Uses host side C-extensions that greatly simplifies host code
- Driver API requires explicit resource management and more verbose syntax
- CUDA C is built on top of Driver API
 - Exposure of Driver API is for historical reasons
 - **No reason to start new development using Driver API**
- Don't confuse the two when referring to CUDA Documentation
 - `cuFunctionName()` – Driver API
 - `cudaFunctionName()` – Runtime API

CUDA Kernel Launch Syntax

- CUDA kernels are launched by the host using a modified C function call syntax:

```
myKernel<<<dim3 dGrid, dim3 dBlock>>>(...)
```

dim3 is vector type with x, y, and z components (eg. dG.x)

Maximum Values For Each Dimension

		Compute Capability		
		2.x (Fermi)	3.x (Kepler) 5.x (Maxwell)	6.x (Pascal)
Total Threads per Block		1024	1024	1024
Grid Size	dGrid.x	65535	$2^{31}-1$	$2^{31}-1$
	dGrid.y	65535	65535	65535
	dGrid.z	65535	65535	65535
Block Size	dBlock.x	1024	1024	1024
	dBlock.y	1024	1024	1024
	dBlock.z	64	64	64

CUDA Kernels

- Denoted by `__global__` function qualifier
 - Eg. `__global__ void myKernel(float* a)`
- Called from host, executed on device
- A few noteworthy restrictions:
 - No access to host memory (in general!)
 - Must return *void*
 - No static variables
 - No access to host functions

CUDA Syntax – Kernels (I)

Kernels can take arguments just like any C/C++ function

- Pointers to device memory
- Parameters passed by value

```
__global__ void SimpleKernel(float* a, float b)
{
    a[0] = b;
}
```

CUDA Syntax – Kernels (II)

Kernels must be declared (but not necessarily defined) in source/header files before they are called

```
// Kernel declaration
__global__ void kernel(float* a);

int main()
{
    dim3 gridSize, blockSize;
    ...
    kernel<<<gridSize,blockSize>>>(a);
}

__global__ void kernel(float* a)
{
    ...
}
```


CUDA Syntax - Kernels

Kernels have read-only built-in variables:

- `gridDim`: dimensions of the grid
 - Uniform for all threads
- `blockIdx`: unique index of a block within grid
- `blockDim`: dimensions of the block
 - Uniform for all threads
- `threadIdx`: unique index of the thread within the thread block
- Cannot vary the size of blocks or grids during a kernel call

CUDA Syntax - Kernels

Built-in variables are typically used to compute unique thread identifiers

- Map local thread ID to a global array index

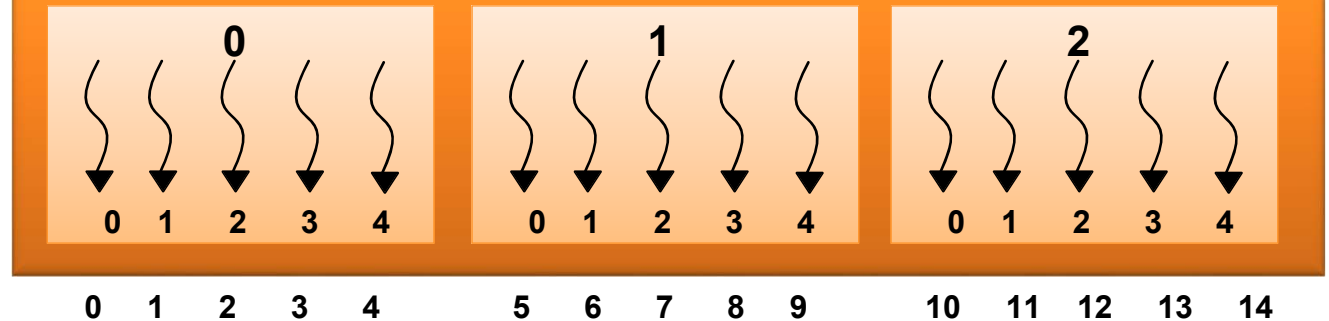
`myKernel<<<3,5>>>(...)`

`blockDim.x = 5`
`gridDim.x = 3`

`blockIdx.x`

`threadIdx.x`

Grid



`blockIdx.x*blockDim.x + threadIdx.x`

CUDA Syntax – Index & Size Calculations

- Global index calculation
 - $idx = blockIdx.x * blockDim.x + threadIdx.x$
- Grid size calculation

$$GridSize = \frac{Size + BlkDim - 1}{BlkDim} \leftarrow \text{Integer Division}$$

- Where
 - Size: Total size of the array
 - BlkDim: Size of the block (max 1024)
 - GridSize: Number of blocks in the grid

CUDA Syntax – Thread Identifiers

Result for each kernel launched with the following execution configuration:

`MyKernel<<<3,4>>>(a);`

```
__global__ void MyKernel(int* a)
{
    int idx = blockIdx.x*blockDim.x+threadIdx.x;
    a[idx] = 7;
}

__global__ void MyKernel(int* a)
{
    int idx = blockIdx.x*blockDim.x+threadIdx.x;
    a[idx] = blockIdx.x;
}

__global__ void MyKernel(int* a)
{
    int idx = blockIdx.x*blockDim.x+threadIdx.x;
    a[idx] = threadIdx.x;
}
```

a: 7 7 7 7 7 7 7 7 7 7 7 7

a: 0 0 0 0 1 1 1 1 2 2 2 2

a: 0 1 2 3 0 1 2 3 0 1 2 3

CUDA Syntax - Kernels

- All C operators are supported
 - eg. +, *, /, ^, >, >>
- All functions from the standard math library
 - eg. sinf(), cosf(), ceilf(), fabsf()
- Control flow statements too!
 - eg. if(), while(), for()

CUDA Kernel C++ Support

- Supported
 - Classes
 - Including inheritance and virtual functions
 - Need to add `__device__` qualifiers to member functions!
 - Templates
 - C++11 features including auto and lambda functions
- Not supported
 - C++ Standard Library
 - Run time type information (RTTI)
 - Exception handling
 - Classes with virtual functions are not binary compatible between host and device

User-defined Device Functions

- Can write/call your own device functions

- `__device__ float myDeviceFunction()`
- Device functions cannot be called by host

```
__device__ float myDeviceFunction(int i)
{
    ...
}

__global__ void myKernel(float* a)
{
    int idx = blockIdx.x*blockDim.x+threadIdx.x;
    a[idx] = myDeviceFunction(idx);
}
```

- Functions declared with both `__device__` and `__host__` will be compiled for both the CPU and GPU

CUDA Syntax - Memory Management

- Typically, host code manages device memory:
 - `cudaMalloc(void** pointer, size_t nbytes)`
 - `cudaMemset(void* pointer, int value, size_t count)`
 - `cudaFree(void* pointer)`

```
// Memory allocation example  
  
int n = 1024;  
int nBytes = 1024*sizeof(int);  
int* a = 0;  
cudaMalloc((void**)&a, nbytes);  
cudaMemset( a, 0, nbytes);  
cudaFree(a);
```

CUDA Syntax – Memory Spaces

- Host and device have separate memory spaces
 - For discrete GPUs data is moved between them via PCIe/NVLink bus
- Pointers are just addresses
 - Can't tell from the pointer value whether the address is on device or host
 - Must exercise caution when dereferencing pointers
 - Dereferencing host pointers on device likely crashed, and vice versa



CUDA Syntax – Data Transfers

- Host code manages data transfers to and from the device:
 - `cudaMemcpy(void* dst, void* src, size_t nbytes, enum cudaMemcpyKind direction);`
 - Direction is one of:
 - `cudaMemcpyHostToDevice`
 - `cudaMemcpyDeviceToHost`
 - `cudaMemcpyDeviceToDevice`
 - `cudaMemcpyHostToHost`
 - `cudaMemcpyDefault`
 - Blocking call - returns once copy is complete
 - Waits for all outstanding CUDA calls to complete before starting transfer
 - With `cudaMemcpyDefault`, runtime determines which way to copy data

CUDA Syntax - Synchronization

- Kernel launches are asynchronous
 - Control returns to CPU immediately
 - Kernel starts executing once all outstanding CUDA calls are complete
- `cudaMemcpy()` is synchronous
 - Blocks until copy is complete
 - Copy starts once all outstanding CUDA calls are complete
- `cudaDeviceSynchronize()`
 - Blocks until all outstanding CUDA calls are complete

```
cudaMemcpy(..., cudaMemcpyHostToDevice);  
  
// Data is on the GPU at this point  
  
MyKernel<<<...>>>(...);  
  
// Kernel is launched but  
// not necessarily complete  
  
cudaMemcpy(..., cudaMemcpyDeviceToHost);  
// CPU waits until kernel is complete  
// and then transfers data  
  
// Data is on the CPU at this point
```

CUDA Syntax – Error Management

- Host code manages errors
- Most CUDA function calls return `cudaError_t`
 - Enumeration type
 - `cudaSuccess` (value 0) indicates no errors
- `char* cudaGetErrorString(cudaError_t err)`
 - Returns a string describing the error condition

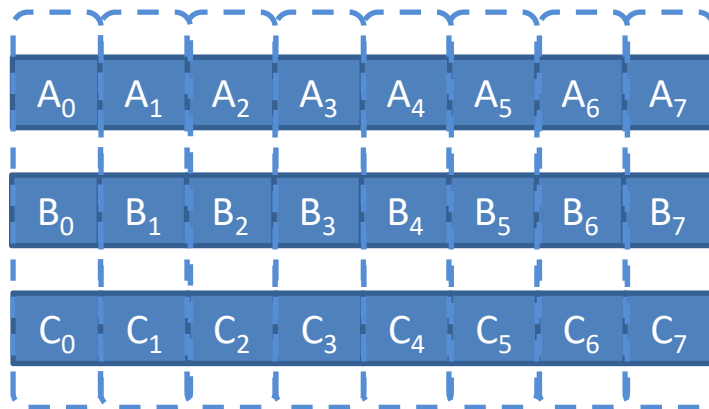
```
cudaError_t e;  
e = cudaMemcpy(...);  
if(e)  
    printf("Error: %s\n", cudaGetErrorString(err));
```

CUDA Syntax – Error Management

- Kernel launches have no return value!
- `cudaError_t cudaGetLastError()`
 - Returns error code for last CUDA runtime function (including kernel launches)
 - Resets global error state to `cudaSuccess`
 - In case of multiple errors, only the last one is reported
 - For kernels:
 - Asynchronous, must call `cudaDeviceSynchronize()` first, then `cudaGetLastError()`

```
MyKernel<<< ... >>> (...);  
  
cudaDeviceSynchronize();  
e = cudaGetLastError();
```

Putting It All Together



One Thread per Output

$$C_x = A_x + B_x \quad \text{Operation}$$

```
__global__  
void VectorAddKernel( float* a,  
                      float* b,  
                      float* c)  
{  
    int idx = threadIdx.x  
            + blockIdx.x * blockDim.x;  
  
    c[idx] = a[idx] + b[idx];  
}
```

This kernel assumes that the size of the array fits evenly into the block size.

What happens if does not?

Vector Add – Host Code

```
void VectorAdd(float* aH, float* bH, float* cH, int N)
{
    float* aD, *bD, *cD;
    int N_BYTES = N * sizeof(float);
    dim3 blockSize, gridSize;

    cudaMalloc((void**)&aD, N_BYTES);
    cudaMalloc((void**)&bD, N_BYTES);
    cudaMalloc((void**)&cD, N_BYTES);

    cudaMemcpy(aD, aH, N_BYTES, cudaMemcpyHostToDevice);
    cudaMemcpy(bD, bH, N_BYTES, cudaMemcpyHostToDevice);

    blockSize.x = 512;
    gridSize.x = N / blockSize.x;
    VectorAddKernel<<<gridSize, blockSize>>>>(aD, bD, cD);

    cudaMemcpy(cH, cD, N_BYTES, cudaMemcpyDeviceToHost);
}
```

This code assumes N
is a multiple of 512

Allocate memory on
GPU

Transfer input
arrays to GPU

Launch kernel

Transfer output
array to CPU

CUDA Syntax – Unified Memory

- Instead of explicitly declaring memory for the host and the device, use managed memory

```
cudaMallocManaged(void **devPtr, size_t size)
```

```
// Memory allocation example  
int n = 1024;  
int nBytes = 1024*sizeof(int);  
int* a = 0;  
cudaMallocManaged((void**)&a, nbytes);  
cudaFree(a);
```

Putting It All Together... Again!

```
int* a, *b, *c;  
int N_BYTES = 2 * sizeof(int);
```

```
cudaMallocManaged((void**)&a, N_BYTES);  
cudaMallocManaged((void**)&b, N_BYTES);  
cudaMallocManaged((void**)&c, N_BYTES);
```



Allocate managed
memory

```
a[0] = 5; b[0] = 7;  
a[1] = 3; b[1] = 4;
```



Initializing memory from
host

```
VectorAddKernel<<<1,2>>>(a, b, c);  
cudaDeviceSynchronize();
```



Launch kernel and
synchronize device

```
printf("%d %d\n", c[0], c[1]);
```


Summary

GPUs are data-parallel architectures

CUDA provides a heterogeneous compute model:

- Host:
 - Memory management (usually)
 - Data transfers
 - Data-parallel kernel launches on device as a grid of thread blocks
 - Error management
- Device (GPU):
 - Executes data-parallel kernels in threads
 - Implemented in C/C++ with a few important extensions, and a few restrictions

Unified memory simplifies transfers and memory management



Acceleware CUDA Training

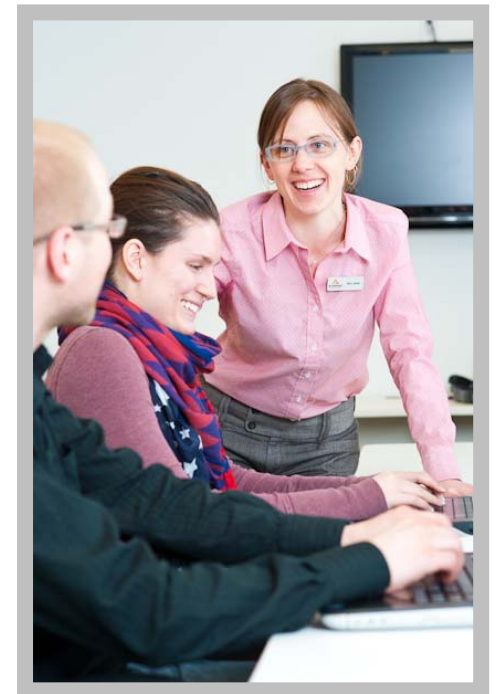
Scheduled CUDA Courses (also available online)

- June 13 – 16: Calgary, Alberta
 - 35% Discount using code: **AXECUDAGTC17**
- September 12 – 15: Calgary, Alberta
- December 5 – 8: Calgary, Alberta

Private training courses

- Courses held onsite at your company
- Delivered anywhere in the world

<http://acceleware.com/cuda-training>



Questions?

Visit us at booth #520

Acceleware Ltd.

Tel: +1 403.249.9099

Email: services@acceleware.com

CUDA Blog: <http://acceleware.com/blog>

Website: <http://acceleware.com>



Chris Mason

chris.mason@acceleware.com

